



# 384.195 Robot Learning Deep Reinforcement Learning Part 2

Dongheui Lee, Yashuai Yan




# Overview

- Summary of value-based RL methods
- Policy-based RL methods
  - Policy Approximation
  - Policy Gradient Theorem
  - REINFORCE
- Actor-Critic RL methods
  - Advantage Actor Critic (A2C)
  - GAE
  - Proximal Policy Optimization (PPO)

# Markov Decision Processes

- Markov decision processes:
  - Set of states  $S$
  - Set of actions  $A$
  - Transitions  $P(s'|s,a)$
  - Rewards  $R(s,a,s')$  (and discount  $\gamma$ )
- MDP quantities so far:
  - Policy = Choice of action for each state
  - Utility = sum of (discounted) rewards

$$V([s_0, s_1, s_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$



# MDP Algorithms

- Value Iteration: Compute optimal values

$$V_{k+1}(s) \leftarrow R(s) + \gamma \max_a \sum_{s'} P(s'|s, a) V_k(s')$$

- Policy Evaluation: Compute values for a particular policy

$$V^\pi(s) \leftarrow R(s) + \gamma \sum_{s'} P(s'|s, \pi(s)) V^\pi(s')$$

- Policy Exaction: use your values to turn into a policy


$$\pi^*(s) = \operatorname{argmax}_a \left[ R(s) + \gamma \sum_{s'} P(s'|s, a) V(s') \right] \quad \begin{array}{l} \text{actions that maximize} \\ \text{Value-Fct} \end{array}$$

- Policy Improvement:

$$\pi_{i+1}(s) = \operatorname{argmax}_a \left[ R(s) + \gamma \sum_{s'} P(s'|s, a) V^{\pi_i}(s') \right] \quad \text{Iteration}$$

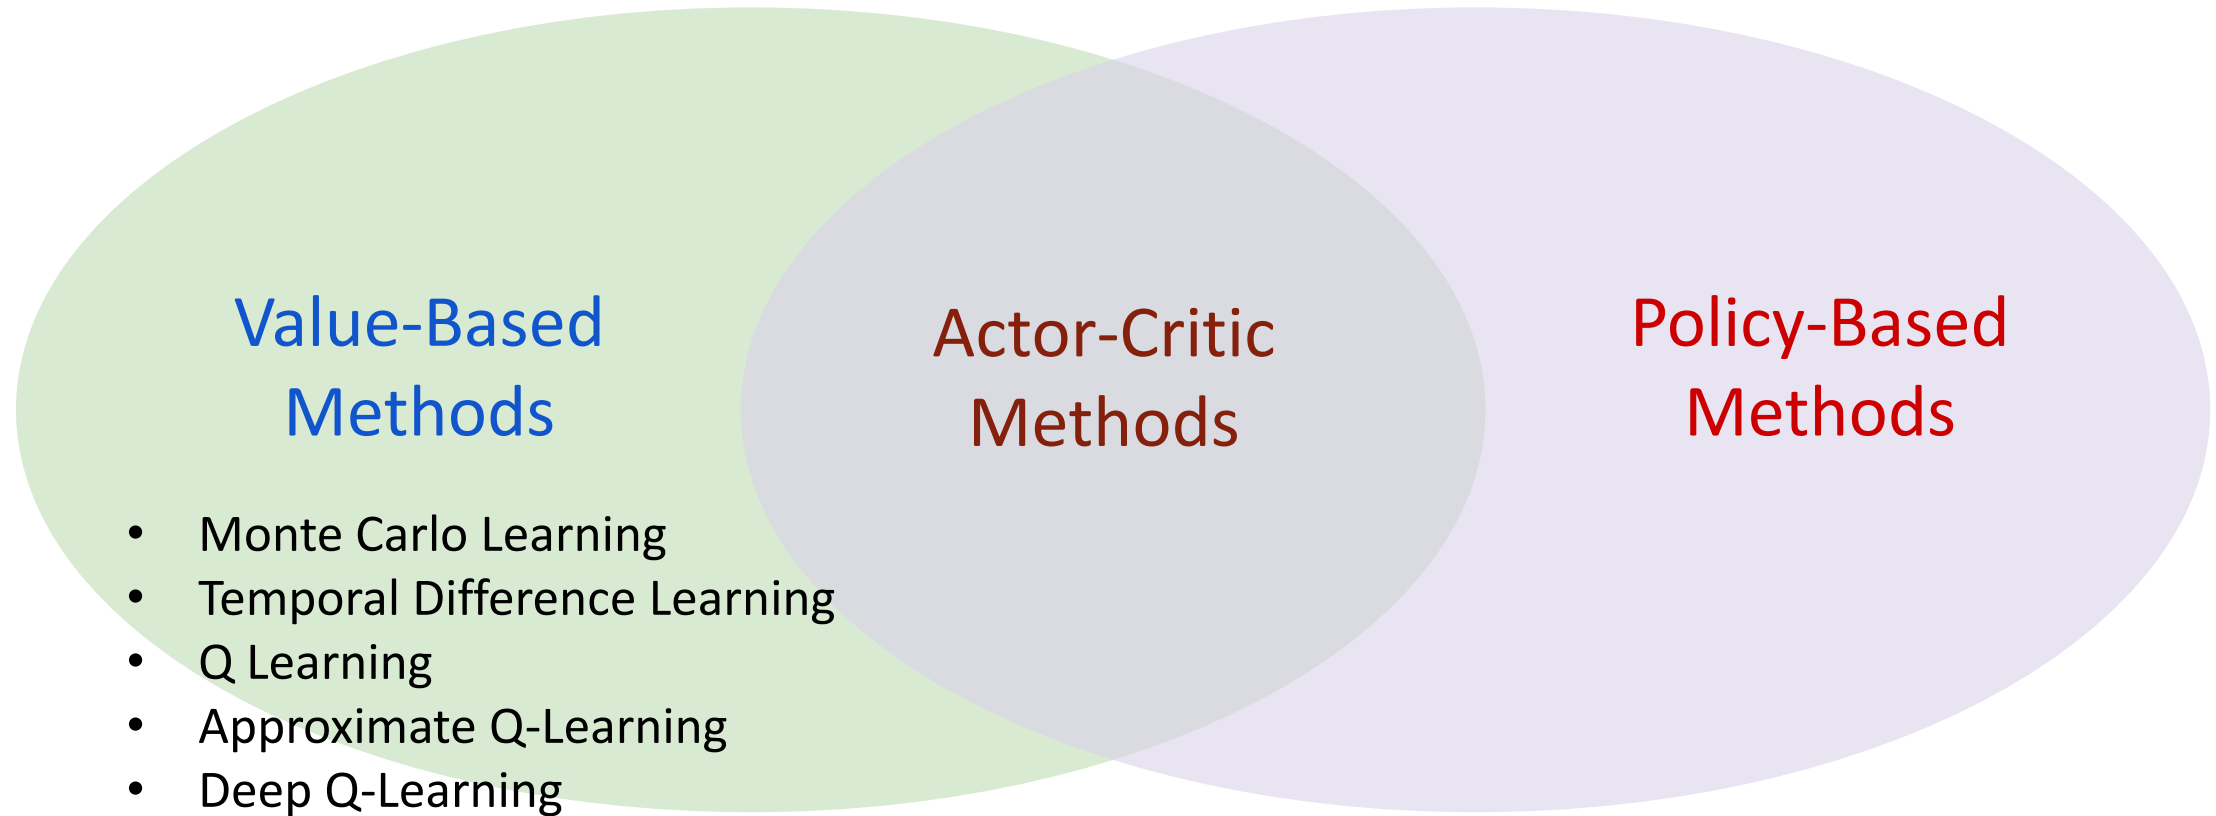
# Reinforcement Learning

- RL involves making a series of optimal actions, it is considered a **sequential decision problem** and can be modeled using Markov Decision Process.
  - A set of states  $s \in S$
  - A set of actions (per state)  $A$
  - A model  $P(s'|s,a)$
  - A reward function  $R(s,a,s')$  (and discount  $\gamma$ )
  - Still looking for a policy  $\pi(s)$
- **Unknown model of transition  $P(s'|s,a)$  and reward  $R(s,a,s')$** 
  - **Sample Backup:** the value of a state is updated from experiences collected from the interaction with the environment.



$$V(s_1) \leftarrow f(r(s_1, a_1) + \gamma V(s_3))$$

# RL methods



# Monte Carlo Learning

- This is called **direct evaluation** or **Monte Carlo on Policy Evaluation**.
- Goal: **learn the (state) values** for each state under  $\pi$
- Idea: **Average together observed sample values**
- It eventually computes the correct average values, using just sample transitions (**unbiased estimator, but high variance**)
- It must **wait until the end of episode** for any update. It takes long to learn.
- It **wastes information about state connections**. Each state must be learned separately.

# Temporal Difference Value Learning

- learn from every experience of a transition  $(s, a, s', r)$ !

- Learn  $V(s)$  like Bellman Update

$$V^\pi(s) \leftarrow R(s) + \gamma \sum_{s'} P(s'|s, \pi(s)) V^\pi(s')$$

- Temporal difference learning of values


- Move values toward the value of whatever successor occurs: running average

Sample of  $V(s)$ :  $sample = R(s, \pi(s), s') + \gamma V^\pi(s')$

Update to  $V(s)$ :  $V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + (\alpha)sample$

Same update:  $V^\pi(s) \leftarrow V^\pi(s) + \alpha(sample - V^\pi(s))$

- However, if we want to turn values into a (new) policy, we're stuck:





# Q-Learning

- Q-Learning: sample-based Q-value iteration

$$Q_{k+1}(s, a) \leftarrow R(s) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q_k(s', a')$$

- Learn  $Q(s, a)$  values as you go

- Receive a sample  $(s, a, s', r)$
- Consider your old estimate:  $Q(s, a)$
- Consider your new sample estimate of  $Q(s, a)$  value:

$$sample = R(s, a, s') + \gamma \max_{a'} Q(s', a') \quad \leftarrow \text{find value for next best state } a'$$

- Incorporate the new estimate into a running average:

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + (\alpha) [sample]$$

- **off-policy learning:** Q-learning converges to optimal policy -- even if you're acting suboptimally!

# From Tabular Q-Learning to Approximate Q-Learning

- Basic Q-Learning keeps a table of all q-values. In realistic situations, it may not be feasible. Too many states or continuous state space.

- Approximate Q-Learning

- Q-Learning with linear feature functions

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha [\text{difference}] \quad \text{Exact Q's}$$

$$w_i \leftarrow w_i + \alpha [\text{difference}] f_i(s, a) \quad \text{Approximate Q's}$$

- Deep Q Network

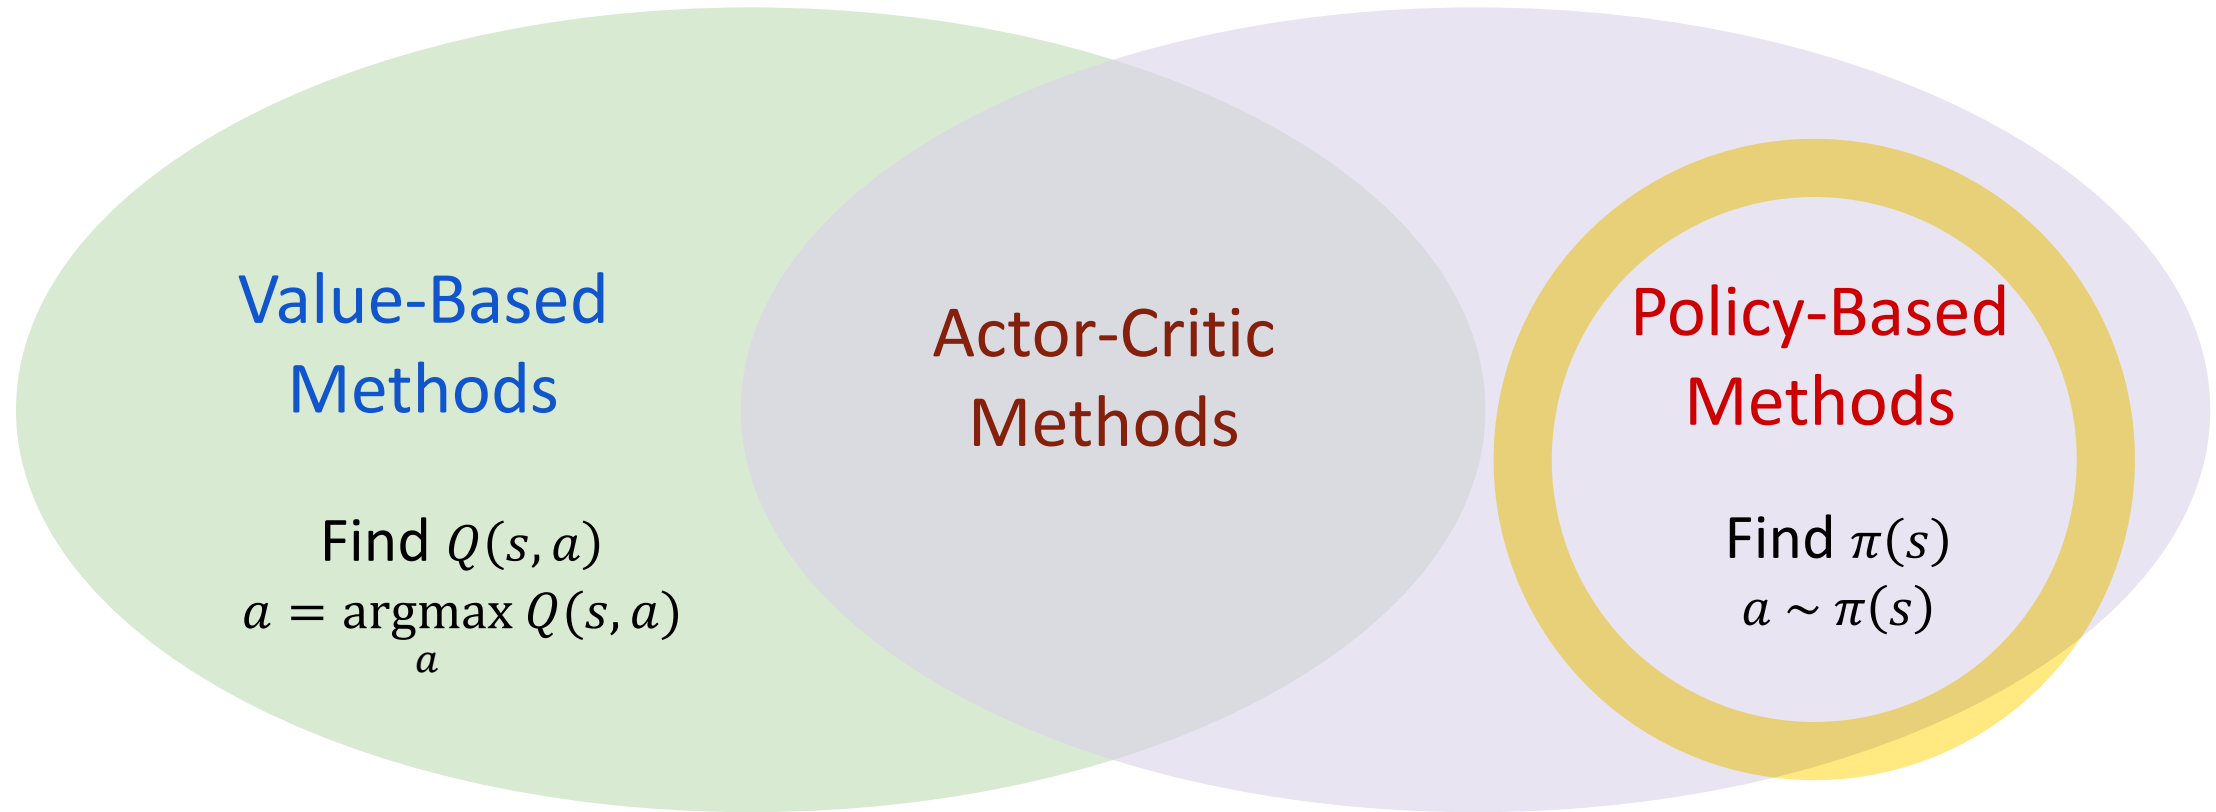


State  $s$

DQN

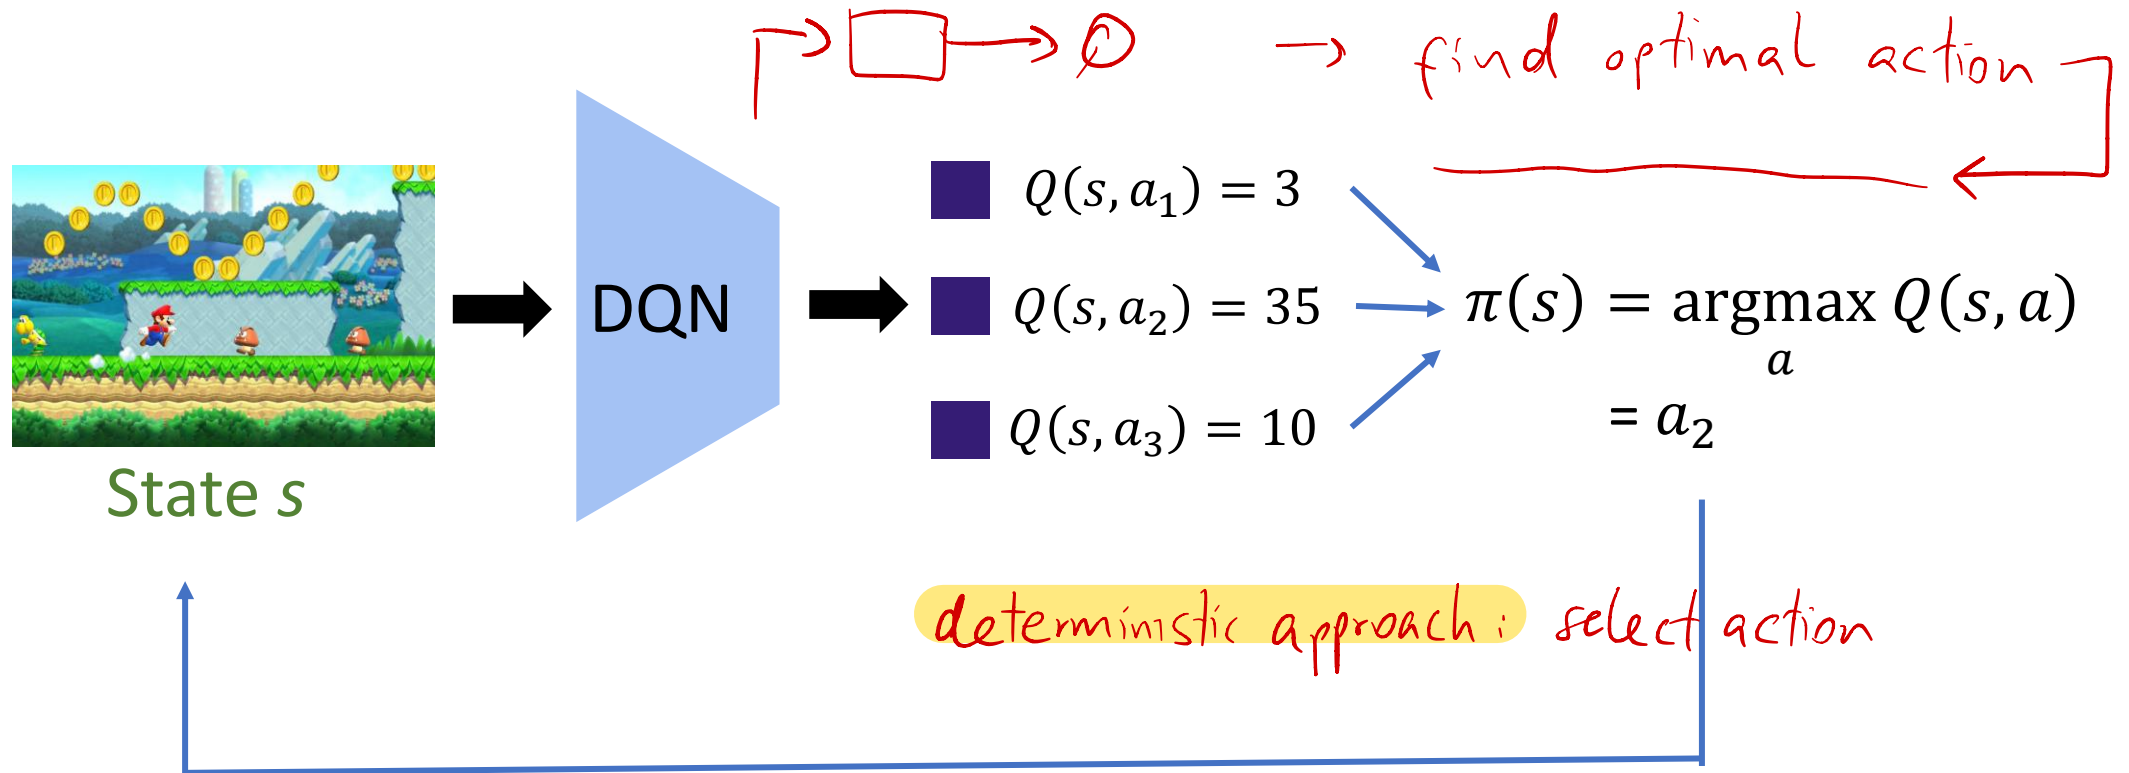
$$\begin{aligned} \blacksquare &\longrightarrow \hat{Q}(s, \text{"left"}; \mathbf{w}) \\ \blacksquare &\longrightarrow \hat{Q}(s, \text{"right"}; \mathbf{w}) \\ \blacksquare &\longrightarrow \hat{Q}(s, \text{"up"}; \mathbf{w}) \end{aligned}$$

# RL methods



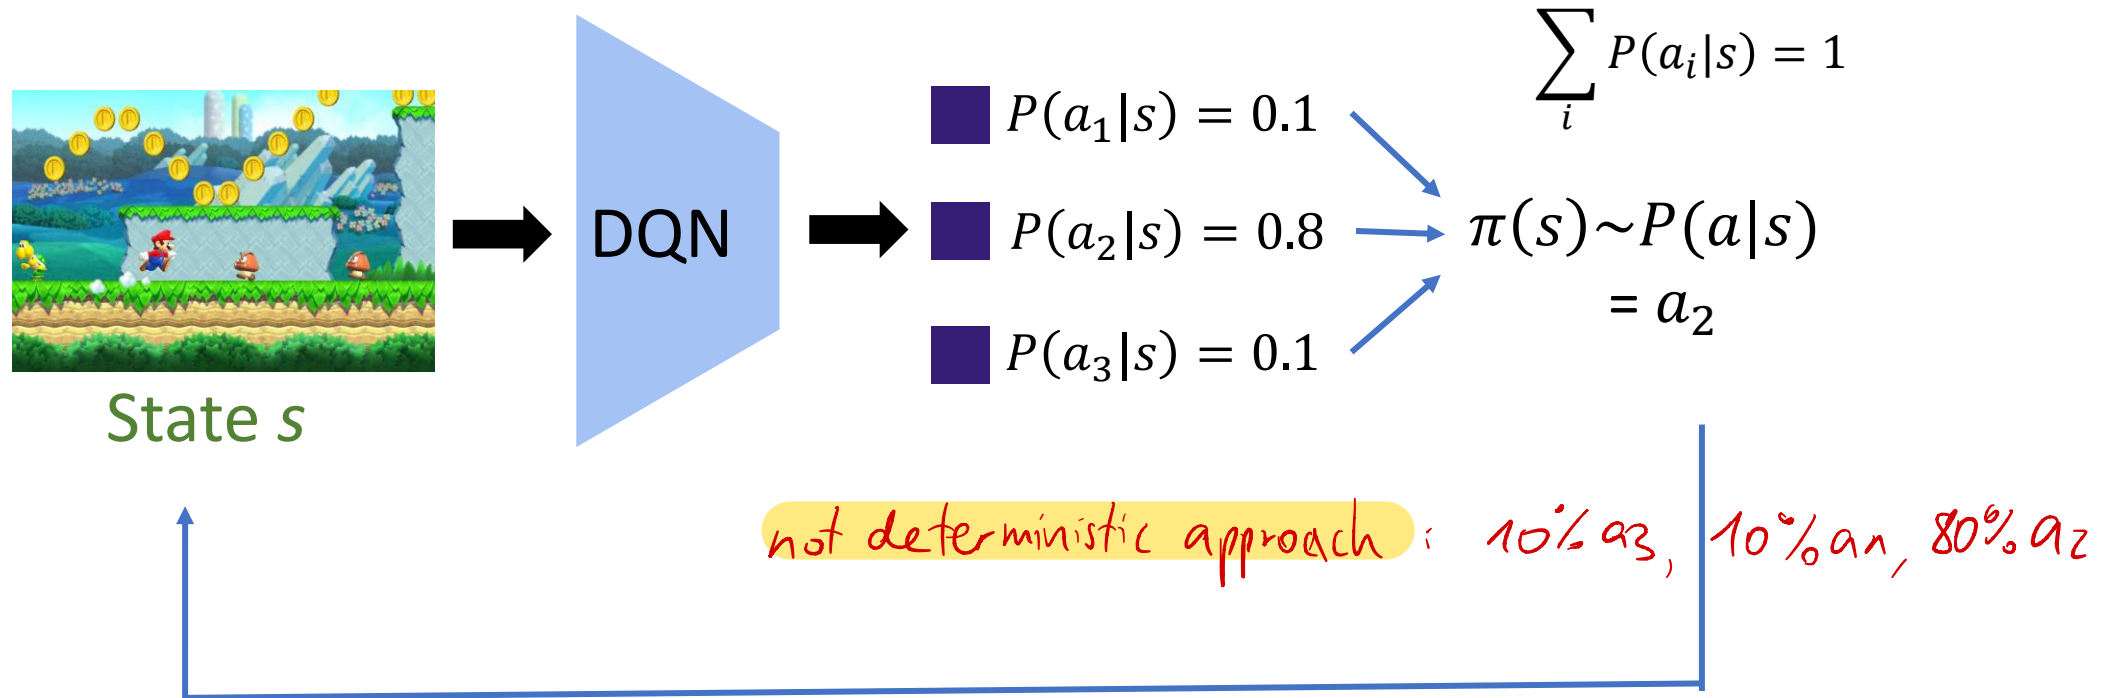
# Deep Q Networks (DQN)

- Use NN to learn Q-function and then use to infer the optimal policy



# Policy Gradient

- DQN: Approximate Q-function and use to infer the optimal policy  $\pi(s)$
- Policy Gradient: Directly optimize the policy  $\pi(s)$




What are the advantages of this formulation?




# Discrete vs Continuous Action Spaces

- Discrete action space: which direction should I move?




# Discrete vs Continuous Action Spaces

- Discrete action space: which direction should I move?
- Continuous action space: how fast should I move? *eg. Gaussian*



# Policy Gradient

- Policy Gradient: Enables modeling of continuous action space



## Downsides of value based reinforcement learning (e.g., DQN)

- $V(s)$  does not prescribe actions
  - Need dynamics model and compute 1 bellman back-up.
- $Q(s, a)$  needs to be able to efficiently solve  $\operatorname{argmax}_a Q(s, a)$ 
  - Complexity: Challenge for continuous & high-dimensional action spaces
- Flexibility
  - Policy is deterministically computed from the Q function by maximizing the reward → cannot learn stochastic policies

To address these, consider a new class of RL training algorithms:  
**Policy gradient methods**

- Often  $\pi$  can be simpler than  $Q(s, a)$  or  $V(s)$

# Advantages of Policy-based RL

- Better convergence properties
- Efficient in high-dimensional or continuous action spaces
- Can learn stochastic policies



# Stochastic Policy

Probability to take action  $a$   
in state  $s$ .



- Consider **stochastic**  $\pi_\theta(a|s) = \pi(a|s; \theta) = P(a|s; \theta)$  **parametrized by  $\theta$** .
- Finitely many **discrete actions**

$$\textbf{Softmax } \pi_\theta(a|s) = \frac{\exp(h(s,a;\theta))}{\sum_{a'} \exp(h(s,a';\theta))}$$

(  $\frac{a}{\sum_{a'} a'}$   $\frac{\text{value for action } a}{\text{all actions}}$  )

where  $h(s, a; \theta)$  might be

**linear**  $h(s, a; \theta) = \sum_i \theta_i f_i(s, a)$

or **non-linear**  $h(s, a; \theta) = \text{NeuralNet}(s, a; \theta)$

- Continuous actions:**

$$\textbf{Gaussian } \pi_\theta(a|s) = N(a|\mu(s; \theta), \Sigma(s; \theta))$$

# Intuition of Policy Gradient RL


## Case study – Autonomous Driving Vehicle

- Agent: Vehicle
- Rewards: distance traveled

State



Camera  
GPS  
Lidar  
...



Action




Steering wheel angle  
Acceleration  
Brake

# Intuition of Policy Gradient RL

## Case study – Autonomous Driving Vehicle

### Policy Gradient Training algorithm

- Initialize the agent
- Run a policy until termination
- Record all states, actions, rewards
- Decrease probability of actions that resulted in low reward → actions near to crash
- Increase probability of actions that resulted in high rewards → actions far away from crash



# How to improve the policy?

perf. measure for recorded policy  $\pi(\theta)$   
↓

- Assume an episodic case. We can define the performance measure  $J(\theta)$  as the value of the start state of the episode.

How to update  $\theta$ ? **policy gradient ascent.**  $\theta \leftarrow \theta + \alpha \nabla J(\theta)$

$$\theta \leftarrow \theta + \alpha \cdot \frac{\partial V(s; \theta)}{\partial \theta}$$

learning rate      policy gradient

How to get performance?

- Challenge: The performance depends on both action selection and the distribution of states (in which those selections are made). Both are affected by the policy parameter.
- How can we estimate the performance gradient wrt the policy parameter when the gradient depends on the unknown effect of policy changes on the state distribution?

# Policy Gradient Theorem

**Policy gradient:** Derivative of  $V(\mathbf{s}; \theta)$  w.r.t.  $\theta$ .

$$\begin{aligned} \bullet \quad \frac{\partial V(\mathbf{s}; \theta)}{\partial \theta} &= \frac{\partial \sum_a \pi(a | \mathbf{s}; \theta) \cdot Q_\pi(\mathbf{s}, a)}{\partial \theta} \\ &= \sum_a \frac{\partial \pi(a | \mathbf{s}; \theta) \cdot Q_\pi(\mathbf{s}, a)}{\partial \theta} \quad \leftarrow \text{no } \theta \text{ parameter} \\ &= \sum_a \frac{\partial \pi(a | \mathbf{s}; \theta)}{\partial \theta} \cdot Q_\pi(\mathbf{s}, a) \quad \int \cdot \frac{\pi(a, s)}{\pi(a, s)} \\ &= \sum_a \pi(a | \mathbf{s}; \theta) \boxed{\frac{\partial \log \pi(a | \mathbf{s}; \theta)}{\partial \theta}} \cdot Q_\pi(\mathbf{s}, a) \\ &= \mathbb{E}_{A \sim \pi(\cdot | \mathbf{s}; \theta)} \left[ \frac{\partial \log \pi(A | \mathbf{s}; \theta)}{\partial \theta} \cdot Q_\pi(\mathbf{s}, A) \right] \end{aligned}$$

(gradient of sum is sum of gradient)

(multiplz by one)

(by property of the gradient of log)

$$\frac{\partial \log \pi(\theta)}{\partial \theta} = \frac{1}{\pi(\theta)} \cdot \frac{\partial \pi(\theta)}{\partial \theta}$$

(by definition of the expectation)



# Policy Update using Policy Gradient Estimate

- **Policy update**  $\theta \leftarrow \theta + \alpha \cdot \frac{\partial V(\mathbf{s}; \theta)}{\partial \theta}$  gradient ascent

- **Policy Gradient:** Derivative of  $V(\mathbf{s}; \theta)$  w.r.t.  $\theta$ .

$$\frac{\partial V(\mathbf{s}; \theta)}{\partial \theta} = \mathbb{E}_{\mathbf{A} \sim \pi(\cdot | \mathbf{s}; \theta)} \left[ \frac{\partial \log \pi(\mathbf{A} | \mathbf{s}; \theta)}{\partial \theta} \cdot Q_{\pi}(\mathbf{s}, \mathbf{A}) \right]$$

computing expectation directly is usually infeasible.

How to calculate policy gradient?

- **Monte-Carlo Estimation:** estimate expectation from random samples.

# Monte Carlo Policy Gradient (REINFORCE)

REINFORCE( $s_0, \pi_\theta$ )

Initialize  $\pi_\theta$  to anything

Loop forever (for each episode)

Generate episode  $s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_T, a_T, r_T$  with  $\pi_\theta$

Loop for each step of the episode  $n = 0, 1, \dots, T$

$\mathcal{O}$ -fct.  $\rightarrow G_n \leftarrow \sum_{t=0}^{T-n} \gamma^t r_{n+t}$  Using return as an unbiased sample of  $Q^\pi(s, a)$

Update policy:  $\theta \leftarrow \theta + \alpha \gamma^n G_n \nabla \log \pi_\theta(a_n | s_n)$

Return  $\pi_\theta$

- Update parameters by stochastic gradient ascent
- Using policy gradient theorem

# Quiz 1



1. What is the fundamental idea behind policy gradient methods in RL?
  - A. Maximizing immediate rewards
  - B. Minimizing action entropy
  - ☒ C. Directly optimizing the policy
  - D. Learning Q-values
  
2. In policy gradient methods, what is the objective function that is typically maximized to improve the policy?
  - A. Q-values
  - B. Maximum likelihood of actions
  - ☒ C. Expected cumulative reward
  - D. Value function approximation

# Quiz

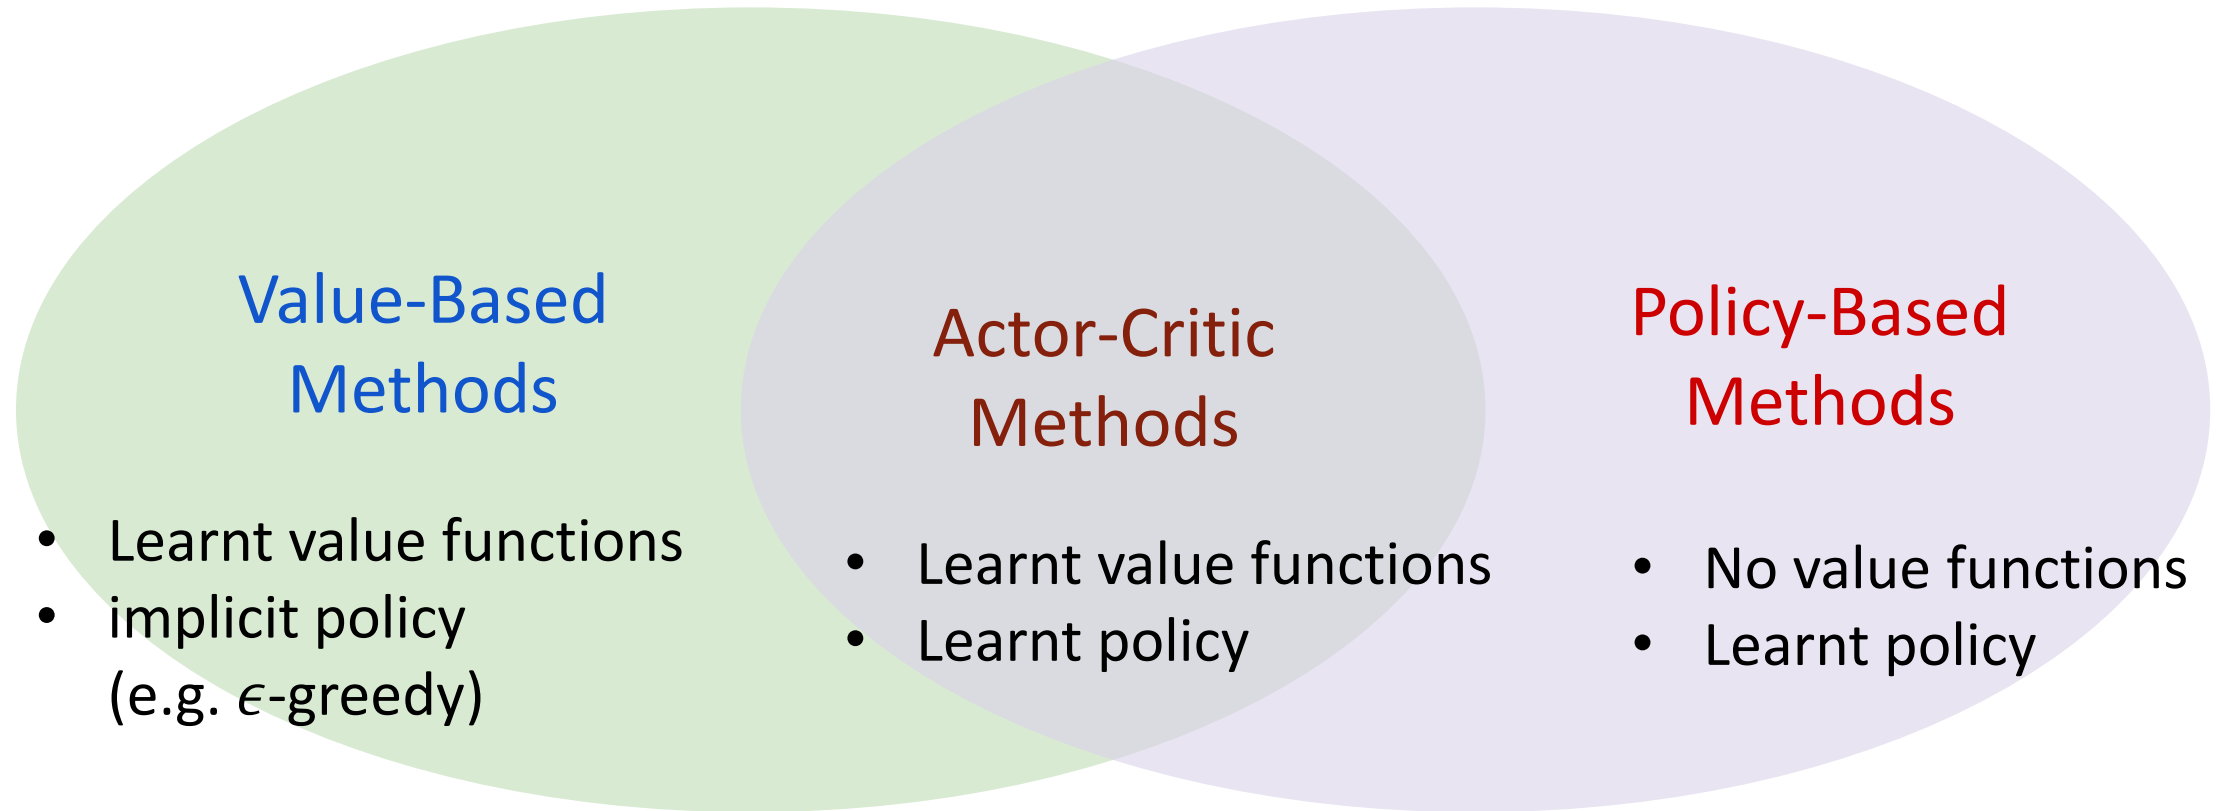
3. How is the policy typically represented in policy gradient methods?

- A. As a fixed set of rules
- B. Using a decision tree
- ☒ C. As a parameterized probability distribution
- D. With a lookup table

4. How do policy gradient methods balance exploration and exploitation?

- A. By always selecting the action with highest probability
- ☒ B. By encouraging stochastic policy
- C. By using epsilon-greedy exploration
- D. By minimizing entropy of the policy

# RL methods





# References

- Sutton and Barto, Reinforcement Learning: An Introduction, Chapter 13